

PEAS: Agent AI chatbot perceives **ENVIRONMENT** (percept; chat interface, third-party plugins, user) through **SENSORS** (text input, chat history, context) → agent function (maps from percept histories to actions) → acts through **ACTUATORS** (text output, image output, API). Based on available info (not nec complete), **Rational agent** choose actions that max **PERFORMANCE MEASURE** (safety, speed, correctness): Best for whom, Maximise what, Unintended effects, Costs.

- **Fully (Partially) observable:** Sensors give agent access to complete state of environment relevant to choice of action.
- **Deterministic:** Next state of environment completely determined by current state and action executed by agent.
- (vs **Stochastic:** some randomness involved eg players choose their move randomly in rock paper scissors)

- **Strategic:** If envt also depend on actions of **other agents**, then is also strategic, unless other actions r predictable, eg dumb
- **Episodic:** Agent's experience divided into atomic episodes (perceive then act), choice of action depends on current episode itself (**Sequential**): Current decision can affect all future decisions, affected by previous sequence of actions.

- **Static (Dynamic):** Environment unchanged when agent is deliberating. (**Semi-dynamic:** Environment does not change with time, but agent's performance score does eg timer).
- **Discrete (Continuous):** Limited no. of distinct, clearly defined percepts & actions.

- **Single (Multi) agent:** Any object(s) that maximise(s) performance measure (value depends on agent's behavior)

Agent is **completely specified by agent function** (environment → sensors → agent function → actuators → environment)

- **Simple reflex agent:** Acts based on current percept, and ignores percept history. Uses condition-action (if-then) rules.

- **Model-based:** Tracks how world changes (transition model) & info from percept; Simulates envt. Condition-action (if-then) rules

- **Goal-based:** Tracks state of the world and what it will be like if it does action A. Picks action that brings agent closer to **goal**.

- **Utility-based:** Also tracks state of world and what it will be like if it does action A. **Utility function** to assign score to any percept sequence (how happy it will be in such a state); internalise performance measure, check alignment.

- **Learning:** Performance element to select external actions; Critic to give feedback on performance; Learning element makes improvement; Problem generator suggests actions

SEARCH PROB (predictable/no opponent → Uninformed / Informed Search): States. Initial State. Goal State(s)/Test. Actions. Transition Model. Action cost function.

— — — **SEARCH ALGORITHMS** — — —
(path matters for (un)informed, adversarial search)

Goal-based, fully observable, deterministic, static, discrete.
⇒ Solution: A fixed sequence of actions → form a path to a goal

Order of node expansion - which states are prioritised

Tree search: Frontier (data structure) defines the search; where it is rare for two paths to reach the same state

Graph search: Remember all **visited** nodes in a hashmap
create **frontier** create **visited** // expands less states
insert initial state to **frontier** and **visited**

```
while frontier is not empty:  
    state = frontier.pop()           // may skip states  
    for action in actions(state):    with less cost  
        next state = transition(state, action)  
        if next state in visited: continue  
        if next state is goal: return solution  
        frontier.add(next state)  
        visited.add(state)  
    return failure
```

UNINFORMED SEARCH (BFS,uniform-cost,DFS,DLS,IDS)

Blind search, no clue how close a state is to the goal(s).

Breadth-First Search (queue) (layer by layer, FIFO) Special case of uniform-cost search (equal step costs). $f(n) = \text{depth of node}$. Save space for BFS while preserving completeness: Apply goal-test when PUSHING a successor state to the frontier; **check if next state is goal BEFORE pushing to frontier queue**

```
create frontier : queue           create visited  
insert initial state to frontier  
while frontier is not empty:  
    state = frontier.pop()           // goal test on the  
    if state is goal: return solution      node popped  
    if state in visited: continue     visited.add(state)  
    for action in actions(state):  
        next state = transition(state, action)  
        // if next state is goal: return solution  
        frontier.add(next state)
```

return failure

Uniform-Cost Search (PRIORITY queue (path cost))

- **A*** with $h(n)=0$: Expand the least cost unexpanded node.
- **Best-First Search** with $f(n) = \text{path cost}$ (FROM INITIAL STATE)

Depth-First Search (stack, LIFO) Can optimise w backtracking search, where only 1 successor generated at a time, O(m) space

Depth-Limited Search (DLS) DFS with depth limit t.

Iterative Deepening Search (IDS) DLS from 0...∞

- IDS NOT always faster than DFS for time complexity. If each state has only a single successor and goal node at depth n, then IDS will run in $O(n^2)$ while DFS will run in $O(n)$.

INFORMED SEARCH (Best-first Search) domain info to guide Best-First Search (priority queue $f(n)$): Eval function $f(n)$ for how good a state is. Expand most desirable unexpanded node.

- **GREEDY Best-First Search** (A^* with $g(n) = 0$) $f(n) = h(n)$

Expand node with lowest $h(n)$, appears closest (low cost) to goal.

- **A* Search** $f(n) = g(n) + h(n)$ $g(n) = \text{cost so far to } n$.

$h(n) = \text{est cost from } n \text{ to goal}$. $f(n) = \text{estimated total cost of path}$.

- $h(n)$ never overestimates cost to reach a goal; conservative.

- If $h(n)$ is admissible → A^* using TREE search is optimal.

Dominance If $h_1(n) \geq h_2(n)$ for all n , then h_2 dominates h_1 and is closer to the true cost $h^*(n)$, incurs less search cost (on average). h_2 is better for search if admissible.

Consistent Heuristics (Is admissible) (satisfies Δ inequality)

- Successor n' of n generated by action a : $h(n') \leq c(n, a, n') + h(n')$ and $h(G) = 0$. $f(n) = \text{non-decreasing along any path}$.

- If $h(n)$ is consistent → A^* using graph search is optimal.

- Consistent heuristic is admissible, but converse not true.

* NO admissibility and consistency in local, adversarial search

Search	Informed, Uninformed	Local
Agent	Goal-based	Goal &/ or Utility-
State space	Low to moderate	Very large. Games.
Time constraint	A solution / No solution	Good enough soln
Solution	Search path (usually)	State

LOCAL SEARCH States, Initial state, Goal test (optional)

- **Successor function:** Possible states from a state

- Objective functions: Evaluate value / goodness of a state
Space landscape: Local max, Global max, Shoulder (same values, but progress possible). *Can still solve shortest path prob* (state=random path; successor=add/remove subpaths)

Hill-Climbing / Steepest Ascent / Greedy Local Search

Find local maxima by traveling to **neighbouring states** with the **highest value**. If node has higher value than successors, then return it. It is a local maximum, but may not be a global max.

Simulated Annealing (hill climbing, but allow bad moves)

Randomly pick successor/action. If the successor has higher value, recurse. Else recurse with probability (=curr. next, T) of choosing a bad move that exponentially decreases. *Escapes local maxima by allowing bad moves occasionally*. If T decrease slowly enough, then finds a global optimum with high probability.

ADVERSARIAL SEARCH AND GAMES (unpredictable)

Fully observable, deterministic, discrete, terminal states exist (infinite run), two-player zero-sum (win-win), turn-taking.

- **States, Initial state, Terminal states** (where game ends)

- **Actions:** Action(s) gives a set of legal moves in s.

- **Transition, Utility function** (output the value of a state)

- **Minimax** (recursive DFS: choose highest minimax value)

- Assume the MIN player plays optimally, trying to minimise player's value. Not optimal against a non-optimal MIN player.

- Pick a move that **maximises player's utility**

Alpha-beta Pruning (prune useless branches)

- Will not change the decision

```
def alpha_beta(state):           [minimax: not highlighted]  
    v = max_value(state, -∞, ∞)  
    return action in successors(state) with value v  
def max_value(state, α, β):     MAX player choose MAX  
    if is_cutoff(state): return eval(state)  
    if is_terminal(state): return utility(state)  
    v = -∞  
    for action, next_state in successors(state):  
        v = max(v, min_value(next_state))  
        α = max(α, v)           // α = highest value for MAX  
        if v >= β: return v // stop if v >= β  
    return v  
def min_value(state, α, β):      β = best choice for MIN  
    if is_cutoff(state): return eval(state)  
    if is_terminal(state): return utility(state)  
    v = ∞  
    for action, next_state in successors(state):  
        v = min(v, max_value(next_state))  
        β = min(β, v)           // β = lowest value for MAX  
        if v <= α: return v // stop if v <= α  
    return v
```

Minimax with cutoff at a certain depth (refer: strikethrough)

- For large/infinite game trees: Cutoff, Use evaluation function

- Replace is_terminal / utility with is_cutoff / eval

- is_cutoff return true if state is terminal, exceed time/depth limit

- eval return utility for terminal, heuristic value for non-terminal

- Evaluation (heuristic) function: Estimate how good a state is.

MACHINE LEARNING: SUPERVISED LEARNING Observe input-output pairs, learn to map input to output. Learn from being given right answers $X \rightarrow Y$. Each eg has correct answer.

- Assume y is generated by a true mapping function $f: x \rightarrow y$. Use a learning algorithm to find a hypothesis $h: x \rightarrow y$ (from Hypothesis class H) s.t. $h \approx f$, given a training set (set may contain errors) $\{(x_1, f(x_1)) \dots (x_n, f(x_n))\}$.

Performance: True measure of hypothesis h (Good if $h \approx f$) is how well it handles inputs it has not seen yet, ie the **test set**.

WEAKLY / SEMI-SUPERVISED LEARNING Correct answer given, but not precise, eg image has face, but dk exact location.

UNSUPERVISED LEARNING Agent learns patterns without labels/explicit feedback. **No correct answers**. Eg **Clustering**

REINFORCEMENT LEARNING Agent learns from reinforcements (rewards, punishments), eg game, maze.

REGRESSION (output is continuous value) → Measure error

Eg Predict housing prices

Predicted $\hat{y}_i = h(x_i)$; true $y_i = f(x_i)$ for N samples $\{(x_1, f(x_1)) \dots (x_N, f(x_N))\}$

• **(Mean) Squared Error**

$$MSE = \frac{1}{N} \sum_{i=1}^N (\hat{y}_i - y_i)^2$$

$$MAE = \frac{1}{N} \sum_{i=1}^N |\hat{y}_i - y_i|$$

CLASSIFICATION → Measure **Correctness and Accuracy**:

$$\text{Accuracy} = \frac{1}{N} \sum_{i=1}^N \mathbb{1}_{\hat{y}_i = y_i}$$

Range is 0 to 1. If prediction $\hat{y}_i = y_i$ (true label). Take mean.

CLASSIFICATION (discrete output) → CONFUSION MATRIX:

Actual Label		Accuracy = $\frac{TP+TN}{TP+FN+FP+TN}$	
		Cancer	Benign
Predicted	Cancer	2	1
	Benign	3	4
		TP	FP
		FN	TN
		$P = \frac{TP}{TP+FP}$	$R = \frac{TP}{TP+FN}$
		$F1 = \frac{2P+R}{P+R}$	$F1 = \sqrt{\frac{PR}{P+R}}$

Recall	TP / (TP+FN)	Precision	TP / (TP+FP)
How many relevant items are selected?	How many positive instances can be recalled (predicted)?	How many selected items are relevant? How precise were the positive predicted instances?	
Maximise if FN type II is very dangerous, eg Cancer (not music) prediction		Maximise this if FP type I error is very costly, eg Email spam, Satellite launch date prediction.	

DECISION TREE FALSE that can always find ≥1 DT to perfectly label every eg in training set, no matter amount of data, attributes. - Data may not be consistent. If we train a DT that determines if a student passes/fails. Can't cleanly split using the training set gender as only attribute. **Attributes**: Binary, Discrete, Binned/Discretised continuous valued (<10, 10-20, 20-40...) **Expressiveness** DT can express any function of input attribute

- Function maps a vector of attribute values to a single output value (a decision eg T/F).
- Internal nodes are tests, leaf specifies value to be returned.
- Trivially, there is a consistent decision tree for ANY TRAINING SET, but probably **will not generalise (prefer a more compact tree)** to new examples.
- Size of Hypothesis class: For Bool functions with n Bool attributes, #distinct decision trees = $\#$ Bool functions = $\#$ distinct truth tables with 2^n rows (each row outputs T/F) → 2^{2^n}

GREEDY, TOP-DOWN, RECURSIVE algorithm for DTL

```
def DTI(examples, attributes, default):  
    if examples is empty: return default  
    if examples have the same classification:  
        return classification  
    if attributes is empty:  
        return classification with highest  
        choose best attribute  
    if attributes is empty:  
        return mode(examples)  
    return mode(examples) define this  
    best = choose attribute(attributes, examples)  
    tree = a new decision tree with root best  
    for each value  $v_i$  of best:  
        examplesi = (rows in examples with best =  $v_i$ )  
        subtree = DTI(examplesi, attributes - best, mode(examplesi))  
        add a branch to tree with label  $v_i$  and subtree subtree
```

- If remaining egs are all True or False, return that value.
- If a mix, then choose best attribute to split and recurse.
- No egs left, return most common value frm node's parent's egs
- No attributes left to split, return most common value of curr egs

Decision Tree Learning with INFORMATION GAIN

Entropy (randomness): $I(P(v_1) \dots P(v_n)) = - \sum_{i=1}^n P(v_i) \log_2 P(v_i)$

- Entropy = 1 if 50:50 eg #true = #false $q = \text{prob of true}$

- Bool variables: $B(q) = -(q \log_2 q + (1-q) \log_2 (1-q))$

- Dataset with p positive and n negative examples:

$$I(\frac{p}{p+n}, \frac{n}{p+n}) = - \frac{p}{p+n} \log_2 \frac{p}{p+n} - \frac{n}{p+n} \log_2 \frac{n}{p+n}$$

Entropy of output of entire set is $B(\frac{p}{p+n})$.

A chosen attribute A with v distinct values divides training set E into subsets $E_1 \dots E_v$ according to their values.

$$\text{remainder}(A) = \sum_{k=1}^v \frac{p_k + n_k}{p+n} I(\frac{p_k}{p_k + n_k}, \frac{n_k}{p_k + n_k})$$

Each subset E_k has p_k positive examples and n_k negative

$$\text{examples}: \text{Remainder}(A) = \sum_{k=1}^v \frac{p_k + n_k}{p_n} B(\frac{p_k}{p_k + n_k})$$

Information Gain (highest = impt attrib) / Reduction in Entropy: $IG(A)$ (in bits) = $I(\frac{p}{p+n}, \frac{n}{p+n}) - \text{remainder}(A)$

= Entropy of this node - Entropy of children nodes

- **Continuous-valued Attributes:** Define a discrete-valued input attribute to partition values into a discrete set of intervals.

- **Missing values:** Assign most common val of A at node n to eg with missing data for A & at node n. Or filter by most common val amongst egs w same output. Or assign each possible value of A some prob, split egs with missing data. Or drop attribute/row.

b: branching factor | d: depth | m: maximum depth | C*: path cost of optimal solution | ε: minimum edge cost | G: goal state

SEARCH	TIME: #nodes generated or expanded	SPACE: maximum #nodes in memory	COMPLETE: Solution / failure if there is	OPTIMAL: Always find least cost
BFS	$1 + b + b^2 + \dots + b^d = O(b^d)$	$O(b^d)$. Worst case: expand last child in a branch.	Complete if b is finite.	Yes if same step cost everywhere.
Uniform-cost	$O(b^{C*/\epsilon})$. Estimated depth = $C*/\epsilon$	$O(b^{C*/\epsilon})$.	Yes if $\epsilon > 0$ and finite C^* .	Yes if $\epsilon > 0$. $\epsilon=0$ may cause zero cost cycle
DFS	$O(b^m)$.	$O(bm)$.	No, when infinite depth / back&forth loop.	No.
DLS	$1 + b + b^2 \dots b^l = O(b^l)$. $l = \text{depth limit}$.	$O(b^l)$ if used with DFS.	No if l is not the depth of optimal solution.	No if used with DFS.
IDS overhead	$b^0 + (b^0 \cdot b^1) \dots + (b^0 \dots + b^d) = (d+1)b^0 + (d+1)b^1 + \dots + (1)b^d = O(b^d)$.	$O(b^d)$ if used with DFS.	Yes.	Yes, if same step cost everywhere.
Greedy Best-First	$O(b^m)$. Good heuristic $h(n)$ gives improve.	$O(b^m)$, keep all nodes in memory.	No, can go into loops.	No, heuristic function can be wrong.
A*	$O(b^m)$. Good heuristic $h(n)$ gives improve.	$O(b^m)$, keep all nodes in memory.	Yes, tracks path so no back and forth. Unless infinite nodes with $f \leq f(G)$.	Depends on heuristics.
Minimax	$O(b^m)$.	$O(bm)$. Depth first exploration (then backtrack)	Yes, if the tree is finite.	Yes, against optimal opponent.
α-β Pruning	$O(b^{m^2})$ with perfect ordering.			

